



# Decision Trees

## An introduction to regression and classification trees

July 5, 2018

Aaron Lowther





- 
- Most notes are taken from Hastie et al. (2009)  
<https://web.stanford.edu/~hastie/ElemStatLearn/>

## Overview...


- Two specific types of *decision trees*:
  - Regression Trees: For a continuous response,
  - Classification Trees: Discrete response.
- Iteratively partition the covariate space into a series of (hyper) rectangles.
- Use a simple function in each rectangle to predict a response.
- We are going to assume we have data,

$$\{(y_t, x_{t,1}, \dots, x_{t,P}) \text{ for } t = 1, \dots, T\},$$

for  $y_t \in \mathbb{R}$  (for now).


# Partitioning

## The covariate space



# Partitioning


## The covariate space



But not like this!

# Partitioning

## The covariate space



But not like this!

We have a grey area, it's hard to explain...

# Partitioning

## Recursive binary partitioning

---




- Instead of arbitrary partitioning
- We stick to recursive binary partitioning
- The advantages are:
  - The partitions are easy to explain
  - There is a simple tree representation of the partitions

# Partitioning



## Recursive binary partitioning: An example

---







# Partitioning



## Recursive binary partitioning: An example

---





# Partitioning



## Recursive binary partitioning: An example



---



# Partitioning





## Recursive binary partitioning: An example



# Partitioning



## Recursive binary partitioning: An example



- We need to determine (automatically):
  - Which covariates to split
  - Where to split them
- With the aim of developing a model of the form,

$$f(\mathbf{x}) = \sum_{m=1}^M c_m \mathbb{I}(\mathbf{x} \in \mathcal{R}_m),$$

- $M$ : Number of regions,
- $\mathcal{R}_m$ : Region  $m$ ,
- $c_m$ : Predicted value of the response for  $\mathbf{x} \in \mathcal{R}_m$
- $\mathcal{R} = \{\mathcal{R}_1, \dots, \mathcal{R}_M\}$

# Making the tree

## Determining the values of $c_m$

- In the linear regression setting we use the least squares criterion,

$$SS(\beta) = \sum_{t=1}^T \left( y_t - \sum_{p=1}^P x_{t,p} \beta_p \right)^2$$

to determine the values of  $\beta$  for the model,

$$y_t = \sum_{p=1}^P x_{t,p} \beta_p + e_t.$$

- We can adopt the same *least squares* criteria here, minimizing

$$\mathcal{RT}(\mathbf{c}; \mathcal{R}) = \sum_{t=1}^T (y_t - f(\mathbf{c}, \mathbf{x}_t))^2$$

# Making the tree

## Determining the values of $c_m$

- We can adopt the same criteria here, minimizing,

$$\mathcal{RT}(\mathbf{c}; \mathcal{R}) = \sum_{t=1}^T (y_t - f(\mathbf{c}, \mathbf{x}_t))^2.$$

- Then the values of  $\mathbf{c}$  that minimise the function,

$$c_m = \frac{1}{N_m} \sum_{t=1}^T y_t \mathbb{I}_{\mathbf{x}_t \in \mathcal{R}_m},$$

are just the mean of the response values in the region (given regions  $\mathcal{R}_m$ ).

# Making the tree

## Determining the regions

---

- Finding  $(c, \mathcal{R})$  that minimises,

$$\mathcal{RM}(c, \mathcal{R}),$$

in general, is very hard.

- Early splits will effect subsequent splits...
- But a greedy approach helps simplify the matter



# Making the tree

## Determining the regions

- To start the search we search each split value for each covariate, finding  $(p, s)$  such that,

$$\min_{p,s} \left[ \min_{c_1} \sum_{x_i \in \mathcal{R}_1(p,s)} (y_i - c_1)^2 + \min_{c_2} \sum_{x_i \in \mathcal{R}_2(p,s)} (y_i - c_2)^2 \right],$$

is obtained.

- $\mathcal{R}_1(p, s) = \{\mathbf{x} | x_p \leq s\}$ ,  $\mathcal{R}_2(p, s) = \{\mathbf{x} | x_p > s\}$
- Apply this splitting approach to subsequent regions.

# Summary so far

---

So now we have a method to build a tree,

- Finding the order to split the covariates,
- Finding the split values for the covariates.

But when do we stop?

# When do we stop?



- 
- We could continue growing the tree providing the reduction in sums of squares is greater than some threshold.
  - The preferred method is to grow a large tree, and then prune it.
  - Cost complexity pruning is used

Let,

- $T_0$ : The largest tree obtained,
- $|T|$ : The number of terminal nodes in  $T$ ,
- $N_m$ : Number of  $\mathbf{x}_t \in \mathcal{R}_m$ ,
- $Q_m(T) = \frac{1}{N_m} \sum_{\mathbf{x}_t \in \mathcal{R}_m} (y_t - \hat{c}_m)^2$

For each  $\alpha$  we find the tree that minimises the cost complexity criteria  $C_\alpha(T)$ ,

$$C_\alpha(T) = \sum_{m=1}^{|T|} N_m Q_m(T) + \alpha |T|,$$

by collapsing the tree by its weakest (internal) node.



When the data is discrete, taking values  $1, 2, \dots, K$  we need to adapt the:

- Slitting criteria.
- Pruning criteria.

For regression we used the *squared error node impurity measure*,

$$Q_m(T) = \frac{1}{N_m} \sum_{\mathbf{x}_t \in \mathcal{R}_m} (y_t - \hat{c}_m).$$

For categorical predictors, three available node impurity measures are:

- Misclassification error,

$$\frac{1}{N_m} \sum_{t \in \mathcal{R}_m} \mathcal{I}((y_t \neq k(m))) = 1 - \hat{p}_{mk(m)}.$$

- Gini Index,

$$\sum_{k \neq k'} \hat{p}_{mk} \hat{m}k' = \sum_{k=1}^K \hat{p}_{mk} (1 - \hat{p}_{mk}).$$

- Cross-entropy/deviance,

$$-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}.$$



- 
- Cross-entropy and Gini index are more sensitive to changes in the node probabilities than the misclassification rate.
  - The cross-entropy or Gini index should be used when growing the tree.
  - Any of the three methods can be used for cost complexity pruning.



---

The textbook covers other issues and extensions that include,

- Categorical Predictors.
- Loss matrix, when misclassifying some observations is more serious in a given class.
- Missing predictor values: A number of approaches, but could use a missing category.
- Linear combinations of splits.



Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning*. Springer Series in Statistics. Springer, second edition.



July 5, 2018  
Aaron Lowther

